

Retrieval Augmented Generation

RAG grounds generation in retrieved passages.

This reduces hallucination on open-domain questions.

Chunking Strategy

Chunk size materially affects retrieval quality.

Small chunks lose context; large chunks dilute embeddings.

Hybrid Search

Combining BM25 with dense vectors outperforms either alone.